MMLU-pro:

⇒ Removed noisy/trivial question from MMLU

⇒ Expanded 4 choices to 10 choices

⇒ Evaluated using CoT (gives model more of a chance)

⇒ Accuracy of models drop by 16% to 33% (not as saturated)

---

GPQA: Graduate-level Google proof Q/A:

⇒ Questions written by 61 phD contractors from upwork

⇒ PhD: >65% accuracy: expert } human
non-phD: 34% accuracy: non-expert } human
GPT-4: 39% accuracy

Question writing
⇓
expert #1 validation
⇓
Question revision
⇓
Expert validation #2
⇑
non-expert validation
⇑
final Q/A set

---

Humanity's Last Exam (HLE): multi-modal, many subjects, multi-choice + short answer

{public HLE/private HLE

Filter by ~~creators~~ Frontier LLMs, multiple stage of review

Benchmarks difficulty: MMLU < MATH < GPQA < HLE

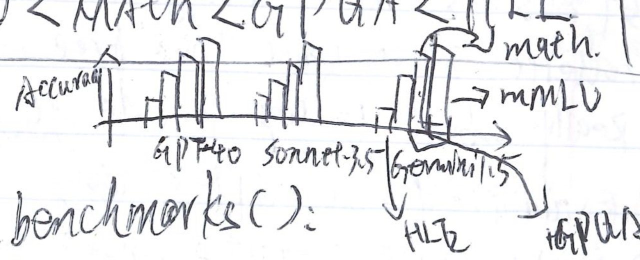


---

~~instruction following~~ #3

Def instruction_following_benchmarks():

{ previous task: structured questions/QA
~~new~~ instruction task: follow the instructions to complete a task

challenge: how to evaluate an open-ended response?

Chatbot Arena: live evaluation on the cloud/air

How it works:

① random persons from the Internet types in prompt.
random uses input prompt

② They get responses from two random(annoymized) models

③ They rate which one is better

④ ELO scores are computed based on the pairwise comparision

⑤ Features: live (not static) inputs, can accomodate new models

AlpacaEval: 805 instructions from various sources

Metric: win rate against GPT-4 preview as judged by GPT-4 preview (potential bias)

WildBench: ⇒ sourced 1024 examples from 1M human-chatbot conversations

⇒ use GPT-4 turbo as a judge with a checklist (e.g. CoT)

⇒ well-correlated (0.95) with chatbot Arena + GPT4 as a judge.

How it works?

History query, LLM-A, Checklist, LLM-B ⇒ pairwise ⇒ WB-reward

{query, History, checklist} + LLM response ⇒ individual ⇒ WB-score