

pre-development safety testing steps:

#1: US safety institute + UK AI safety institute working together

#2: Company give model access to AI institute before model

#3: safe institutes run evaluations and produce a release report to

↓
Hallucination in medical setting makes systems capable and more safe.
↓
1. hallucination itself is bad
2. ~~can~~ understanding/handling hallucination makes system stronger

two aspects of a model to reduce: capability + propensity / more safer
⇒ A system is able to do something, but refuse to do it

⇒ for API models, propensity matters

⇒ For open weight models, capability matters since they can easily fine-tune safety away

Most of benchmarks are for away from real-world use

these benchmarks: are # # quizzing-based prompt: user knows the answer, and try to test the system (like standard exam)

⇒ # asky-based prompt: users don't know the answer, and try to ask the system to get it

{ asky-base prompt is more realistic and produce values for users

e.g. Med HELM: previous medical benchmarks were based on standard exam
121 clinical tasks sourced from 29 clinicians, mixture of private/public data