

Example: OLMO from AI2 2024

- pretraining: wikipedia / corpus / wiki book / ~~get~~ arXiv
- midtraining: o/a / math / coding ds
- post-training: safety / multilingual / precise IP.

⌋ (*)
What are these data?
How are they chosen / processed?

Some models with trained data:

def
⇒ Bert(): (books corpus, wikipedia) even high-quality data
might contain malicious contents

def
⇒ Gpt2-webtext(): webtext / OpenWebTextCorpus

def (*) common-crawl():
↑ Reddit link ↑ de-duplicates text from website

● servers:
two formats:
↓
WARC: raw http responses (HTML) / JSON
WET: converted to text (lossy process) (*) released
↓
CCNet

use Apache Nutch crawler:
run web crawl monthly (starts with some seed urls)
~100 crawls from 2013-2025 (download pages in a queue / add hyperlink)
in 2016, crawl takes 10-12 days on 100 machines
Latest crawl: April 2025
crawls have some overlap, but try to diversify

⇒ def CCNet(): goal: automatic way of constructing large, high-quality datasets for pre-training.
↓ CCNet

includes/contains: De-duplication
Language identification and classification
Quality filtering: keep docs from wiki under a / perplexity 5-gram model

Results: Trained BERT models, CCNet outperforms wiki
CCNet: