

def $T_5(C)$: colossal clean crawled corpus (C)
 

text-to-text transfer transformer model

rule-based filtering: #1 keep lines that end in punctuation and have ≥ 5 words
#2: remove pages with fewer than 3 sentences

#3: Remove page that contains any 'bad' words

#4: Remove page containing 'no code', 'lorem ipsum',

English only #5: Filter out non-English 'term of use', etc.
text using langdetect

Result: 806 GB of text (156 billion tokens)

def $GPT_3(C)$: { common crawl: CC
webText2: webtext with external links
internet-based books corpora.
wiki

Result: 570 GB of 400B tokens

process by common crawl: { trained quality classifier to tell: webtext, wiki, -
fuzzy deduplication of documents: webtext / benchmarks

def thePile(C): curated 22 high-quality domains
composition of Academic, Internet, Prose, Dialogue, Misc.

Result:

0. in reaction to GPT-3

1. 825 GB of text (~275B tokens)

2. pile - CC: common crawl, use WARC

3. PubMed Central: 5 million pages, NIH-funded work

4. Arxiv: preprint since 1991 5. Enron emails: 500k, 150 users from

youtube
title/captions
comments

Enron management, release during Enron investigation.

def: project-gutenberg():

2025: 175K English books

copyright clearances book.

PG-19: books from project-gutenberg before 2019

def books3():

196K books from ~~shadow library~~ Bibliotik.

Contained books from authors

Has been taken down due to copyright infringement/lawsuits

Shadow libraries: Library Genesis (LibGen), Z-library, Sci-Hub

Open trained models on Lib Gen

def ~~exchange~~ stackexchange():

1: started with stackoverflow in 2008, grew to other eqires

2: collection of sites of user-contributed questions/answers

3: Q/A format is close to instruction tuning/real application

4: Data dumps in XML.

def github(): 3.1TB of code

started in 2008, acquired by Microsoft in 2018

repos of code bases / meta data / PRs / commits

Lots of duplicates (copied code, fork)

Using minhash and Jaccard similarity to remove near-duplicates