




Dr. GRPO doesn't do
sed normalization

no length issue, but easily stuck in suboptimal
issue
tuning hyperparameters
solve it ---

● Lecture #17 - Summary

- ①: RL is the key to surpassing human abilities
- ②: if you can measure it, you can optimize it
- ③: policy gradient framework is conceptually clear,
just need baselines to reduce variance
- ④: RL system is more complex than pre-training,
including inference workloads, manage multiple
models

